

Homework 4: Exponential Smoothing

2025-02-01

Question 7.1

Describe a situation or problem from your job, everyday life, current events, etc., for which exponential smoothing would be appropriate. What data would you need? Would you expect the value of α (the first smoothing parameter) to be closer to 0 or 1, and why?

Exponential smoothing could be useful if a brand is dropping a new clothing line and they want to predict how many pieces they'll sell over the first few weeks. Demand is unpredictable, maybe there's a huge surge at launch, then sales slow down, or maybe a celebrity wears it, and sales spike again. Exponential smoothing could help adjust their forecast in real time based on how sales are actually going.

What Data Would Help? Early sales numbers (how fast things sell in the first few hours/days) Pre-order & website traffic (how many people showed interest before launch) Marketing impact (influencer posts, ads, and hype levels) Stock & restocks (if it sells out, does demand stay high or drop off?)

Would α Be Closer to 0 or 1? Closer to 1 (like 0.7 - 0.9) because demand can change fast. If a product is blowing up on social media, the forecast needs to adjust quickly. If sales slow down after launch, the model should pick up on that too. For something more consistent, like a basic hoodie that always sells at a steady rate, α would be lower (closer to 0) to smooth out random spikes. But for more unique or bold pieces, the model needs to react fast to sudden surges in demand.


Question 7.2

Using the 20 years of daily high temperature data for Atlanta (July through October) from Question 6.2 (file temps.txt), build and use an exponential smoothing model to help make a judgment of whether the unofficial end of summer has gotten later over the 20 years. (Part of the point of this assignment is for you to think about how you might use exponential smoothing to answer this question. Feel free to combine it with other models if you'd like to. There's certainly more than one reasonable approach.)

Note: in R, you can use either HoltWinters (simpler to use) or the smooth package's es function (harder to use, but more general). If you use es, the Holt-Winters model uses model="AAM" in the function call (the first and second constants are used "A"dditively, and the third (seasonality) is used "M"ultiplicatively; the documentation doesn't make that clear).


```
#load the temps data
data <- read.table('temps.txt', stringsAsFactors = FALSE, header = TRUE)
library(ggplot2)

data_vec <- as.vector(unlist(data[,2:21]))
#convert vector data to time series data with 123 observations for 20 years.
data_ts <- ts(data_vec, start = 1996, frequency = 123)
plot(data_ts)
```




```
#check if there is trend and seasonality via decompose function
plot(decompose(data_ts))
```

Decomposition of additive time series



```
#model_1 for single exponential smoothing
model_1 <- HoltWinters(data_ts, beta = FALSE, gamma = FALSE)
plot(model_1$fitted)
```

model_1\$fitted




```
#model_2 for double exponential smoothing
model_2 <- HoltWinters(data_ts, gamma = FALSE)

## Warning in HoltWinters(data_ts, gamma = FALSE): optimization difficulties:
## ERROR: ABNORMAL_TERMINATION_IN_LNSRCH
model_2

## Holt-Winters exponential smoothing with trend and without seasonal component.
##
## Call:
## HoltWinters(x = data_ts, gamma = FALSE)
##
## Smoothing parameters:
##   alpha: 0.8445729
##   beta : 0.003720884
##   gamma: FALSE
##
## Coefficients:
##      [,1]
## a 63.2530022
## b -0.0729933
plot(model_2$fitted)
```

model_2\$fitted



```
#model_3 for triple exponential smoothing with seasonal=additive
model_3 <- HoltWinters(data_ts)
model_3
```


```
## Holt-Winters exponential smoothing with trend and additive seasonal component.
##
## Call:
## HoltWinters(x = data_ts)
##
## Smoothing parameters:
##   alpha: 0.6610618
##   beta : 0
##   gamma: 0.6248076
##
## Coefficients:
##               [,1]
## a      71.477236414
## b     -0.004362918
## s1    18.590169842
## s2    17.803098732
## s3    12.204442890
## s4    13.233948865
## s5    12.957258705
## s6    11.525341233
## s7    10.854441534
## s8    10.199632666
## s9     8.694767348
## s10    5.983076192
## s11    3.123493477
```

## s12	4.698228193
## s13	2.730023168
## s14	2.995935818
## s15	1.714600919
## s16	2.486701224
## s17	6.382595268
## s18	5.081837636
## s19	7.571432660
## s20	6.165047647
## s21	9.560458487
## s22	9.700133847
## s23	8.808383245
## s24	8.505505527
## s25	7.406809208
## s26	6.839204571
## s27	6.368261304
## s28	6.382080380
## s29	4.552058253
## s30	6.877476437
## s31	4.823330209
## s32	4.931885957
## s33	7.109879628
## s34	6.178469084
## s35	4.886891317
## s36	3.890547248
## s37	2.148316257
## s38	2.524866001
## s39	3.008098232
## s40	3.041663870
## s41	2.251741386
## s42	0.101091985
## s43	-0.123337548
## s44	-1.445675315
## s45	-1.802768181
## s46	-2.192036338
## s47	-0.180954242
## s48	1.538987281
## s49	5.075394760
## s50	6.740978049
## s51	7.737089782
## s52	8.579515859
## s53	8.408834158
## s54	4.704976718
## s55	1.827215229
## s56	-1.275747384
## s57	1.389899699
## s58	1.376842871
## s59	0.509553410
## s60	1.886439429
## s61	-0.806454923
## s62	5.221873550
## s63	5.383073482
## s64	4.265584552
## s65	3.841481452

s66 -0.231239928
s67 0.542761270
s68 0.780131779
s69 1.096690727
s70 0.690525998
s71 2.301303414
s72 2.965913580
s73 4.393732595
s74 2.744547070
s75 1.035278911
s76 1.170709479
s77 2.796838283
s78 2.000312540
s79 0.007337449
s80 -1.203916069
s81 0.352397232
s82 0.675108103
s83 -3.169643942
s84 -1.913321175
s85 -1.647780450
s86 -5.281261301
s87 -5.126493027
s88 -2.637666754
s89 -2.342133004
s90 -3.281910970
s91 -4.242033198
s92 -2.596010530
s93 -7.821281290
s94 -8.814741200
s95 -8.996689798
s96 -7.835655534
s97 -5.749139155
s98 -5.196182693
s99 -8.623793296
s100 -11.809355220
s101 -13.129428554
s102 -16.095143067
s103 -15.125436350
s104 -13.963606549
s105 -12.953304848
s106 -16.097179844
s107 -15.489223470
s108 -13.680122300
s109 -11.921434142
s110 -12.035411347
s111 -12.837047727
s112 -9.095808127
s113 -5.433029341
s114 -6.800835107
s115 -8.413639598
s116 -10.912409484
s117 -13.553826535
s118 -10.652543677
s119 -12.627298331


```
## s120 -9.906981556  
## s121 -12.668519900  
## s122 -9.805502547  
## s123 -7.775306633
```

```
plot(model_3)
```



```
plot(model_3$fitted)
```

model_3\$fitted



```
#model_4 for triple exponential smoothing with seasonal=multiplicative
model_4 <- HoltWinters(data_ts, seasonal = 'multiplicative')
model_4
```


```
## Holt-Winters exponential smoothing with trend and multiplicative seasonal component.
##
## Call:
## HoltWinters(x = data_ts, seasonal = "multiplicative")
##
## Smoothing parameters:
##   alpha: 0.615003
##   beta : 0
##   gamma: 0.5495256
##
## Coefficients:
##              [,1]
## a      73.679517064
## b     -0.004362918
## s1      1.239022317
## s2      1.234344062
## s3      1.159509551
## s4      1.175247483
## s5      1.171344196
## s6      1.151038408
## s7      1.139383104
## s8      1.130484528
## s9      1.110487514
## s10     1.076242879
## s11     1.041044609
```


s12 1.058139281
s13 1.032496529
s14 1.036257448
s15 1.019348815
s16 1.026754142
s17 1.071170378
s18 1.054819556
s19 1.084397734
s20 1.064605879
s21 1.109827336
s22 1.112670130
s23 1.103970506
s24 1.102771209
s25 1.091264692
s26 1.084518342
s27 1.077914660
s28 1.077696145
s29 1.053788854
s30 1.079454300
s31 1.053481186
s32 1.054023885
s33 1.078221405
s34 1.070145761
s35 1.054891375
s36 1.044587771
s37 1.023285461
s38 1.025836722
s39 1.031075732
s40 1.031419152
s41 1.021827552
s42 0.998177248
s43 0.996049257
s44 0.981570825
s45 0.976510542
s46 0.967977608
s47 0.985788411
s48 1.004748195
s49 1.050965934
s50 1.072515008
s51 1.086532279
s52 1.098357400
s53 1.097158461
s54 1.054827180
s55 1.022866587
s56 0.987259326
s57 1.016923524
s58 1.016604903
s59 1.004320951
s60 1.019102781
s61 0.983848662
s62 1.055888360
s63 1.056122844
s64 1.043478958
s65 1.039475693

s66 0.991019224
s67 1.001437488
s68 1.002221759
s69 1.003949213
s70 0.999566344
s71 1.018636837
s72 1.026490773
s73 1.042507768
s74 1.022500795
s75 1.002503740
s76 1.004560984
s77 1.025536556
s78 1.015357769
s79 0.992176558
s80 0.979377825
s81 0.998058079
s82 1.002553395
s83 0.955429116
s84 0.970970220
s85 0.975543504
s86 0.931515830
s87 0.926764603
s88 0.958565273
s89 0.963250387
s90 0.951644060
s91 0.937362688
s92 0.954257999
s93 0.892485444
s94 0.879537700
s95 0.879946892
s96 0.890633648
s97 0.917134959
s98 0.925991769
s99 0.884247686
s100 0.846648167
s101 0.833696369
s102 0.800001437
s103 0.807934782
s104 0.819343668
s105 0.828571029
s106 0.795608740
s107 0.796609993
s108 0.815503509
s109 0.830111282
s110 0.829086181
s111 0.818367239
s112 0.863958784
s113 0.912057203
s114 0.898308248
s115 0.878723779
s116 0.848971946
s117 0.813891909
s118 0.846821392
s119 0.819121827


```
## s120 0.851036184  
## s121 0.820416491  
## s122 0.851581233  
## s123 0.874038407
```

```
plot(model_4)
```




```
plot(model_4$fitted)
```

model_4\$fitted



```
#calculate mean of each day across the 20 years
model_3_season <- matrix(model_3$fitted[,4], nrow = 123)
avg_daily <- data.frame(rowMeans(model_3_season, n = 19))
ggplot(avg_daily, aes(x = 1:123, y = rowMeans.model_3_season..n...19.)) +
  geom_line() +
  labs(
    x = "Day of the Year (1 to 123)", # x-axis label
    y = "Average Daily Temperature"  # y-axis label
  )
```



```
std <- sd(avg_daily$rowMeans.model_3_season...19.[1:62])
C <- std*0.3
T <- std*5
C
```

```
## [1] 0.9697854
```

```
T
```

```
## [1] 16.16309
```

```
s = list()
avg_daily$s <- 0
#cusum method from hw3 is applied to check change to the end of summer
cusum <- function(C, T, cl){
  for (j in cl){
    for (i in 1:nrow(avg_daily)){
      if (i-1 != 0){
        s <- max(0, avg_daily$s[i-1] + mean(avg_daily[,j]) - avg_daily[i,j] - C)
      }
      else{s <- 0}
      if (s >= T){s}
      else{s <- 0}
      avg_daily$s[[i]] <- s
    }
  }
  return(data.frame(1:123, avg_daily[,j], avg_daily$s)[78:123,])
  # set output start from base level day 78
}
```

```
cusum(C=0.9697854, T=16.16309, cl=1)
```

##	X1.123	avg_daily...j.	avg_daily.s
## 78	78	-0.6562640	0
## 79	79	-0.9865849	0
## 80	80	-1.6968615	0
## 81	81	-1.6688708	0
## 82	82	-1.7134622	0
## 83	83	-3.0786163	0
## 84	84	-3.3870835	0
## 85	85	-1.1096552	0
## 86	86	-2.0995425	0
## 87	87	-1.5440073	0
## 88	88	-1.9147525	0
## 89	89	-2.2592635	0
## 90	90	-2.8837074	0
## 91	91	-6.3260493	0
## 92	92	-8.6200910	0
## 93	93	-9.3828956	0
## 94	94	-9.4880777	0
## 95	95	-6.5151863	0
## 96	96	-7.6117861	0
## 97	97	-9.2529673	0
## 98	98	-13.4282487	0
## 99	99	-16.2629184	0
## 100	100	-14.9832377	0
## 101	101	-15.5134511	0
## 102	102	-13.2152299	0
## 103	103	-12.2627214	0
## 104	104	-12.2140597	0
## 105	105	-11.8562254	0
## 106	106	-10.5370460	0
## 107	107	-9.4835355	0
## 108	108	-7.9795695	0
## 109	109	-5.6027596	0
## 110	110	-7.9832251	0
## 111	111	-9.8836170	0
## 112	112	-10.6011249	0
## 113	113	-7.1261238	0
## 114	114	-7.7851236	0
## 115	115	-11.5537389	0
## 116	116	-11.6623496	0
## 117	117	-10.0716372	0
## 118	118	-7.4072725	0
## 119	119	-8.8583903	0
## 120	120	-7.6735617	0
## 121	121	-4.4656976	0
## 122	122	-1.2369297	0
## 123	123	-0.8479661	0

Has Summer Been Ending Later in Atlanta? A Look at 20 Years of Temperature Data

This analysis explores whether the end of summer in Atlanta has been getting later over the past 20 years (1996–2015) by looking at daily high temperatures from July through October. Using exponential smoothing models and CUSUM (Cumulative Sum Control Chart) analysis, we tried to detect any meaningful shifts in late-summer temperatures.

Breaking Down the Models

To smooth out fluctuations and identify trends, we used different types of Holt-Winters exponential smoothing models:

- **Model 1 Single smoothing:** Only smooths the temperature data, assuming no trend or seasonality.
- **Model 2 Double smoothing:** Incorporates trend but assumes no seasonal variation.
- **Model 3 Triple smoothing (additive seasonality):** Assumes temperature increases or decreases by a fixed amount each year.
- **Triple smoothing (multiplicative seasonality):** Assumes seasonal temperature variations scale proportionally rather than remain constant.

The trend coefficient in our models came out to be close to zero (-0.0044), which suggests that the overall timing of late summer temperatures hasn't shifted much. While there are small variations, the seasonal patterns have remained fairly stable.

Checking for Changes Using CUSUM

To detect whether summer has been ending later, the CUSUM method was applied. This technique accumulates small changes in temperature patterns to identify significant deviations. The key thresholds were determined using the standard deviation of early summer temperatures:

- **Control Limit (C):** 30% of the standard deviation, used to filter out normal fluctuations. $C = 0.9697854$ is the threshold to account for minor noise that don't indicate a real shift.
- **Threshold (T):** 5 times the standard deviation, defining a significant shift. If the cumulative deviation exceeds $T = 16.16309$, it is a signal of a true shift or change in the temperature pattern. If the cumulative sum reaches or exceeds this value, it suggests that a meaningful change in the data has occurred, beyond just random noise.

The results showed no clear sign that summer has been ending later. While there were some fluctuations in later years, they weren't strong enough to be considered a real trend.

Conclusion

Based on these models, there's **no strong evidence that the unofficial end of summer in Atlanta has been shifting later** over the past two decades. The seasonal temperature patterns look pretty stable, and while there may be small changes year to year, they aren't statistically significant.